

# CSC209 Summer 2015 — Software Tools and Systems Programming

[www.cdf.toronto.edu/~csc209h/summer/](http://www.cdf.toronto.edu/~csc209h/summer/)

Week 2 — May 21, 2015

Peter McCormick  
pdm@cs.toronto.edu

Some materials courtesy of Karen Reid

# Labs

Last Name	Room	TA
A-H	BA2270	Daniel Kats
I-M	BA2240	Alexey Khrabrov
N-Z	BA2220	Michael Chiu Pan Zhang

# Asking for help

*“It doesn’t work”*

*“How do I do XYZ?”*

*“I get an error on line 10”*

# Asking for help

*“I tried **X***

*and expected **Y***

*but got **Z** instead.*

*Please help!”*

# Asking for help

*I tried the following code but instead of printing 3.1415 which I expected, it instead prints 3.000000. What am I doing wrong? Please help!*

```
#include <stdio.h>
int main() {
    int pi = 3.1415;
    printf("%f\n", (double) pi);
    return 0;
}
```

# Agenda

- Introduction to the C language
- The memory model of the machine

# The C Problem Language

- C is a high-level language — structured
  - Supports functions, records and some forms of code modularity
  - Not as high-level as Python or Java
- C is a low-level language — machine level access
- C is a small language — relatively simple syntax, with libraries for extensibility
- C does not hold your hand — it assumes that you know what you're and how you want to do it

# The C Problem Language

- **Good:**
  - Efficient
  - Powerful
  - Portable
  - Flexible
- **Bad:**
  - Easy to make errors
  - Obfuscation
  - Weak support for modularization



From Java to C

# Common Syntax between Java and C (1)

- Distinction between *statements* and *expressions*
- Semicolon denotes end of statement
- Whitespace is generally ignored!
- Braces to denote scope:  

```
    { statement1;  
      statement2; ... }
```

# Common Syntax between Java and C (2)

- Binary Expressions:
  - Comparison: ==, !=, <, <=, >, >=
  - Arithmetic: +, -, /, \*, %
  - Boolean logical: && (and), || (or)
  - Bitwise logical: & (and), | (or), ^ (xor)
  - Bitwise shift: << (left), >> (right)

# Common Syntax between Java and C (3)

- Unary expressions:
  - Minus: -
  - Logical negation: ! (not)
  - Bitwise negation/flip: ~ (not)
  - Pre- and post- increment and decrement (with side effect): ++, --
- Ternary conditional: ?:

# Common Syntax between Java and C (4)

- More expressions:
  - Assignment: =
  - Operator assignment: +=, -=, \*=, /=, %=, &=, |=, ^=, <<=, >>=
- Statement
  - `return expression` (or just `return` if void return type)
- Declarations:
  - `int variable;`
  - `short var1, var2;`
  - `double array[10];`

# Common Syntax between Java and C (5)

- Loops
  - `for`, `while`, `do-while`
  - `break` and `continue` statements
- `if` and `if-else`
- `switch` (with `case` and `default`)

# Compiling C

```
$ gcc -Wall -g -o hello hello.c
```

`-Wall`

Include all warnings. Helps you prevent errors.

`-g`

Include debugging symbols. Allows you to debug with gdb

`-o hello`

Produce an executable called `hello`

`hello.c`

The list of source files to compile.

# main function

- Entry point for all programs
- Each shell argument is passed in as a string
- Standard signature: (not quite true)

```
int main(int argc, char *argv[])
```

- Returns an *exit status*: non-0 indicates an error occurred, otherwise 0 for success



# C data types

- Basic types and literals (King: Ch 7)

```
int i = 38;          long el = 38L;  
int hex = 0x2a;      int oct = 033;  
printf("i = %d, el = %ld, hex = %d, oct = %d\n",  
       i, el, hex, oct);
```

```
i = 38, el = 38, hex = 42, oct = 27
```

```
double d1 = 0.3;     double d2 = 3.0;  
double d3 = 6.02e23;  
printf("d1 = %f, d2 = %f, d3 = %e\n", d1, d2, d3)
```

```
d1 = 0.300000, d2 = 3.000000, d3 = 6.020000e+23
```

# C literals and types

Literal	Value	Type
38	38	int
38L	38	long int
0x2a (hex)	42	int
033 (octal)	27	int
38.0	38.0	double
38.0f	38.0	float

# C data types

- Most things in C are ints:
  - Boolean values are ints
    - 0 means false, nonzero means true
  - characters are ints (ASCII code)
    - 'a' == 97, '\n' == 10, '\033' == 033 == 27
  - enumerations are really ints
- signed vs. unsigned types
- char, int, long, ... are just different sizes of integers.

# Mixed Mode Arithmetic

```
double m = 5/6; /* int / int = int */  
printf("Result of 5/6 is %f\n", m);
```

Result of 5/6 is 0.000000

```
double n = (double)5/6; /* double / int = double */  
printf("Result of (double)5/6 is %f\n", n);
```

Result of (double)5/6 is 0.833333

```
double o = 5.0/6; /* double / int = double */  
printf("Result of 5.0/6 is %f\n", o);
```

Result of 5.0/6 is 0.833333

```
int p = 5.0/6; /* double / int = double but then  
               converted to int */  
printf("Result of 5.0/6 is %d\n", p);
```

Result of 5.0/6 is 0

# Data Type Conversion

- The expression on the right side is converted to the type of the variable on the left.

```
char c;  
int i = c;      /* c is converted to int */  
double d = i;  /* i is converted to double */
```

- This is no problem as long as the variable's type is at least as "wide" as the expression.

```
char c = 500; /* compiler warning */  
int k = d;  
printf("c = %c, k = %d\n", c, k);
```

c = , k = 0

# printf and format strings

- `printf(a_string)` will print the given string
- *Variadic*: `printf` can take a variable number of arguments
- Whether it actually does will depend on special *format strings*:
  - **%d** for signed integers: `printf("%d + %d = %d\n", -3, 5, 2)`
  - **%s** for strings: `printf("Hello %s!\n", "CSC209")`
  - **%f** for floating point: `printf("pi ~ %f\n", 3.14f)`
  - **%c** for ASCII character: `printf("C%cC209", 'S');`
  - **%%** to print an actual %: `printf("100%%!\n")`
  - Other modifiers available: look them up with *man 3 printf*

# Boolean values in C

- No builtin `bool` type, nor `true` and `false` values!
- 0 is considered to be *false*, anything else is *true*
- `if (0) { printf("Never run\n"); }`
- `if (-1) { printf("Always run\n"); }`

# Data Type Capacity

- What happens when the following code is executed?

```
char c = 127;  
int d;
```

```
printf("c = %d\n", c);  
c++;
```


```
d = 512 / c;  
printf("c = %d, d = %d\n", c, d);
```



# Memory Model

# Memory Model

- System memory is can be viewed as a sequence of *bytes* (8 bit values)
- Each location in that sequence (and thus its associated value) is assigned a unique *address*
- Each address is just a number:



# Memory Model

A 32 bit address can give a unique address number  
to ~4 billion ( $2^{32}$ ) different bytes

4294967296 bytes

~4294967 thousand bytes

~4295 million bytes

~4 billion bytes

*aka* ~4.29 gigabytes

== 4 gibibytes ( $4 \times 2^{30}$ )

# Memory Model

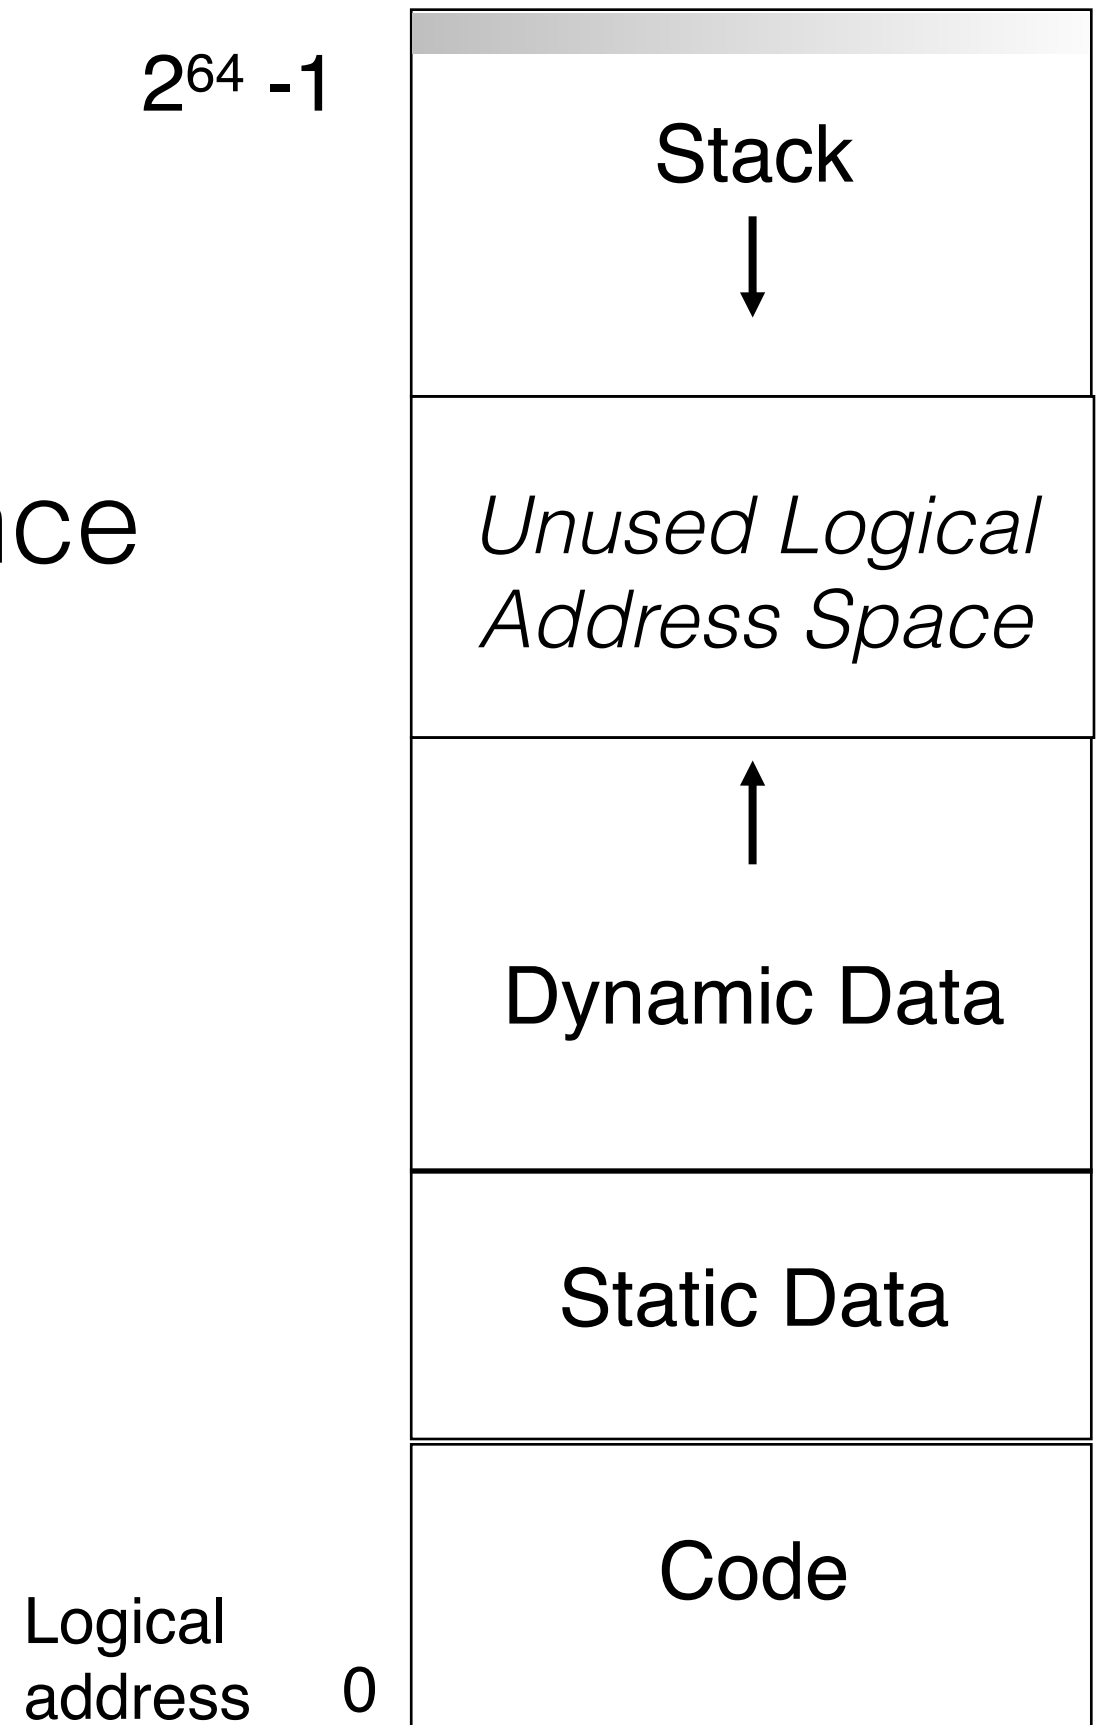
- A 32-bit system can address, and thus is limited to, a maximum of 4GB of system memory (RAM)
- A 64-bit system has a much higher limit (~16 billion GB worth of unique addresses, less usable in practise)
  - The CDF server *Wolf* is a 64-bit machine (with 64GB of physical RAM)
  - This is indicated by the string "x86\_64" in the output of `uname -m`

# Memory Model

- Java and Python hide (shield?) all of this from you
- C does not

# Logical Memory M

- Memory is just a sequence of bytes
- A memory location is identified by an address



# Example

```
main { 0x7fffffffffea9c i  
      f { 0x7fffffffffea7c j  
        0x7fffffffffea6c q  
        0x7fffffffffea68 p  
  
int x = 10;  
int y;  
  
int f(int p, int q) {  
    int j = 5;  
    return p * q + j;  
}  
  
int main() {  
    int i = x;  
    y = f(i, i);  
    return 0;  
}
```

0x601030 y

0x601018 x

10
5
10
10
Unused Logical Address Space
Dynamic Data
???
10
Code

# Arrays

- Arrays in C are a contiguous chunk of memory that contain a list of items of the same type.
- If an array of ints contains 10 ints, then the array is 40 bytes. There is nothing extra.
- In particular, the size of the array is not stored with the array. There is *no* runtime checking.



# Arrays

```
int x[5];  
for (i = 0; i <= 5; i++) {  
    x[i] = i*i;  
}
```

x[0]

x[1]

x[2]

x[3]

x[4]

?

0x88681140

0x88681144

0x88681148

0x8868114c

0x88681150

0x88681154

- No runtime checking of array bounds
- Behaviour of exceeding array bounds is “undefined”
  - program might appear to work
  - program might crash
  - program might do something apparently random

# Next Week

- Assignment 1 will be posted within the next few days
- Lecture: More on C pointers and memory

# Labs

Last Name	Room	TA
A-H	BA2270	Daniel Kats
I-M	BA2240	Alexey Khrabrov
N-Z	BA2220	Michael Chiu Pan Zhang