

CSC209 Summer 2015 — Software Tools and Systems Programming

www.cdf.toronto.edu/~csc209h/summer/

Week 6 — June 18, 2015

Peter McCormick
pdm@cs.toronto.edu

Some materials courtesy of Karen Reid

Announcements

- **Midterm:**
 - Tuesday, June 23 from 7-8pm in UC 266

Announcements

- **Extra Office Hours:**
 - Friday, June 19, 1-3pm in BA3289
 - Monday, June 22, 1-3pm in BA3201

Announcements

- **Assignment 2:**
 - Has been released
 - Due on July 1, 2015 (*less than 2 weeks*)
 - No lecture after this one, before due date

Announcements

- **Labs:**
 - Consolidating down to **BA2210** and **BA2220** only
 - Everyone attending should muster there (there will *not* be TA's in the other two rooms)

Agenda

- Standard library string functions
- Using the C Preprocessor for Sharing Definitions
- Makefiles
- Midterm

*void **

sizeof (void) == 1

... even though *void* is an *incomplete* type,
has no values and you cannot instantiate
any variable to be of type *void* ...


```
void *p = (void *) 0x1000;
```

```
// p + 1 == (void *) 0x1001
```

```
// p + 2 == (void *) 0x1002
```

```
// p + 3 == (void *) 0x1003
```

```
// p + 4 == (void *) 0x1004
```

```
// ...
```

Strings in C

char

char

- The `char` datatype represents an 8-bit integer (is *signed* by default, alternatively can be declared as *unsigned*)
- Small integers can fit in a `char` sized value holder:
 - `char x = 100;`
- Oversized values will be truncated and wrap around:
 - `1 == (unsigned char) 257`
- One ASCII character enclosed in *single quotes* is a `char` literal value:
 - `'A' == 65` since the ASCII code for `A` is decimal 65
 - `'0' == 48` since the ASCII code for `0` is decimal 48
 - `'\0' == 0` uses an escape code
 - `'\n' == 10` (linefeed aka newline)

Strings in C

- A `NULL` (zero) terminated array/sequence of `char` (*byte*) values
 - Typically passed around as a pointer (*`char *`*) to the first character
 - No extra information about length or maximum size
 - Modified *in place* (not necessarily copied)
- *String literals* will include an implicit `NULL` byte
 - `char s[] = "CSC"; s[3] == 0`

strrepr.c:

```
char course1[] = { 67, 83, 67, 50, 48, 57, 0 };  
char course2[] = { 'C', 'S', 'C', '2', '0', '9', '\0' };  
char course3[] = "CSC209";
```

$\forall 0 \leq i \leq 6:$

`course1[i] == course2[i] == course3[i]`

C Standard Library String Functions

- A small but useful set of functions that help you to manipulate C-style strings
 - They require care and attention to detail when using
 - Many traps for young players; source of *many* bugs
- Don't forget to `#include <string.h>` to get the function prototypes

strlen

`strlen` - calculate the
length of a string

```
size_t strlen(const char *s)
```

From the manpage: “*The `strlen()` function calculates the length of the string `s`, excluding the terminating null byte (`'\0'`)..*”

String *Length*
vs
Storage *Size*

strlen.c:

```
char hello[] = { 'c', 's', 'c', '2', '0', '9', '\0' };  
  
printf("hello          = \"%s\"\\n", hello);  
printf("sizeof (hello) = %zu\\n", sizeof (hello));  
printf("strlen(hello)   = %zu\\n", strlen(hello));
```

hello = "CSC209"
sizeof (hello) = 7
strlen(hello) = 6

length + 1 ≤ size

strcmp

```
strcmp("CSC209", "CSC209") == 0
```

```
strcmp("CSCB09", "CSC209") > 0
```

since 'B' > '2' (ASCII 66 vs 50)

```
strcmp("CSC209", "CSC309") < 0
```

since '2' < '3' (ASCII 50 vs 51)

strchr *and* **str***r***chr**

`strchr` & `strrchr` — locate first/last occurrence of a character within a string

```
char *strchr (const char *s, int c)  
char *strrchr(const char *s, int c)
```

From the manpage: “*The `strchr()`/`strrchr()` function returns a pointer to the first/last occurrence of the character `c` in the string `s`.*”

strchr.c

strcat

`strcat` - append (concatenate)
one string onto another

```
char *strcat(char *dest, const char *src)
```

From the manpage: “*The `strcat()` function appends the `src` string to the `dest` string, overwriting the terminating null byte (`'\0'`) at the end of `dest`, and then adds a terminating null byte. ... returns `dest`.*”

strcat.c:

```
#include <stdio.h>
#include <string.h>


int main(int argc, char *argv[])
{
    char s[16];
    int i;

    s[0] = '\0';


    // Concatenate each `argv[i]` onto `s`
    for (i = 1; i < argc; i++) {
        strcat(s, argv[i]);
        strcat(s, "!");
    }

    printf("%s\n", s);
    return 0;
}
```


```
argv[1] = "hello"  
argv[2] = "csc209"  
argv[3] = "fun"
```




```
strcat(s, argv[1]);  
strcat(s, "!");  
  
strcat(s, argv[2]);  
strcat(s, "!");  
  
strcat(s, argv[3]);  
strcat(s, "!");
```




```
strcat(s, "hello");  
strcat(s, "!");  
strcat(s, "CSC209");  
strcat(s, "!");  
strcat(s, "fun");  
strcat(s, "!");
```




→ `strcat(s, "hello");`
`strcat(s, "!");`
`strcat(s, "CSC209");`
`strcat(s, "!");`
`strcat(s, "fun");`
`strcat(s, "!");`




→ `strcat(s, "hello");`
`strcat(s, "!");`
`strcat(s, "CSC209");`
`strcat(s, "!");`
`strcat(s, "fun");`
`strcat(s, "!");`




 `strcat(s, "hello");`
`strcat(s, "!");`
`strcat(s, "CSC209");`
`strcat(s, "!");`
`strcat(s, "fun");`
`strcat(s, "!");`




```
strcat(s, "hello");  
→ strcat(s, "!");  
strcat(s, "CSC209");  
strcat(s, "!");  
  
strcat(s, "fun");  
strcat(s, "!");
```




```
strcat(s, "hello");  
strcat(s, "!");  
→ strcat(s, "CSC209");  
strcat(s, "!");  
  
strcat(s, "fun");  
strcat(s, "!");
```




```
strcat(s, "hello");  
strcat(s, "!");  
→ strcat(s, "CSC209");  
strcat(s, "!");  
strcat(s, "fun");  
strcat(s, "!");
```







```
strcat(s, "hello");  
strcat(s, "!");  
strcat(s, "CSC209");  
→ strcat(s, "!");  
strcat(s, "fun");  
strcat(s, "!");
```




```
strcat(s, "hello");  
strcat(s, "!" );  
strcat(s, "CSC209");  
strcat(s, "!" );  
→ strcat(s, "fun" );  
strcat(s, "!" );
```




```
strcat(s, "hello");  
strcat(s, "!");  
strcat(s, "CSC209");  
strcat(s, "!");  
→ strcat(s, "fun");  
strcat(s, "!");
```



```
strcat(s, "hello");  
strcat(s, "!");  
strcat(s, "CSC209");  
strcat(s, "!");  
→ strcat(s, "fun");  
strcat(s, "!");
```

```
strcat(s, "hello");  
strcat(s, "!");  
strcat(s, "CSC209");  
strcat(s, "!");  
→ strcat(s, "fun");  
strcat(s, "!");
```



s[0]


s[15]

s[16]

```
strcat(s, "hello");  
strcat(s, "!");  
strcat(s, "CSC209");  
strcat(s, "!");  
→ strcat(s, "fun");  
strcat(s, "!");
```

h	e	l	l	o	!	C	S	C	2	0	9	!	f	u	n
s[0]															s[15]

```
strcat(s, "hello");  
strcat(s, "!");  
strcat(s, "CSC209");  
strcat(s, "!");  
strcat(s, "fun");  
→ strcat(s, "!!");  
s[16] = '!';  
s[17] = '\0';
```



Caveat: This was on a
Mac OS X laptop...

Try “*fun12345678*” on
CDF for similar effect...

What was in memory at `s[16]` and `s[17]`,
and what other purpose was it serving?

Regardless, `strcat` went further in memory than ***we*** should have allowed it.

Solution: keep track of how big `dest` will *need* to be, and correctly allocate sufficient space using `malloc`.

This is error prone and
very easy to get wrong.

strcat2.c:

...

```
size_t len = 0;
```

```
// Add up all string lengths
```

```
for (i = 1; i < argc; i++) {
```

```
    len += strlen(argv[i]);
```

```
    // Add +1 for each '!'
```

```
    len++;
```

```
}
```

```
// Another +1 for the NUL terminator
```

```
char *s = (char *) malloc(len + 1);
```

```
s[0] = '\0';
```

...

strcat

- Only use it once you've determined the size of `src` and know that `dest` has sufficient *free* space to accommodate
- *No way of telling `strcat` how large `dest` is (so that it won't accidentally go beyond the end)*

strcpy

strcpy - copy a string

```
char *strcpy(char *dest, const char *src)
```

From the manpage: “*The `strcpy()` function copies the string pointed to by `src`, including the terminating null byte (`'\0'`), to the buffer pointed to by `dest`. The strings may not overlap, and the destination string `dest` *must be large enough* to receive the copy. ... returns `dest`.*”

Whereas `strcat` looks for the NUL terminator of `dest` and then appends the contents of `src` starting there, `strcpy` overwrites from the beginning of `dest`.

```
char s[1024];  
s[0] = '\0';  
  
strcat(s, "Hello ");  
strcat(s, "CSC209!");  
printf("%s\n", s);
```

Hello CSC209!

```
char s[1024];  
s[0] = '\0';  
  
strcpy(s, "Hello ");  
strcpy(s, "CSC209!");  
printf("%s\n", s);
```

CSC209!

Another trap: what happens if `src` is larger than `dest`?

strcpy2.c:

```
char s[16];  
s[0] = '\0';
```

```
strcpy(s, "The quick brown fox jumps over the lazy dog");
```

Introducing **str***n*cpy

strncpy - copy a string

```
char *strncpy(char *dest,  
               const char *src,  
               size_t n)
```

From the manpage: “The `strncpy()` function is similar, except that **at most `n`** bytes of `src` are copied.

Warning: If there is no zero byte among the first `n` bytes of `src`, the string placed in `dest` will not be zero terminated.”

It is your responsibility to
keep track of how many bytes
are free/unused in `dest`.

strncpy.c

You should *always* use `strncpy`,
and *never* use `strcpy` because
it is so outrageously unsafe.

See also `strncmp` and `strncat`
(but they still only deal with `src`
size, not `dest` size.)

strdup

strdup - duplicate a string

```
char *strdup(const char *s)
```

From the manpage: “*The `strdup()` function returns a pointer to a new string which is a duplicate of the string `s`. Memory for the new string is obtained with `malloc(3)`, and can be freed with `free(3)`..*”

It's not unusual for a C standard library function to modify a string argument *in-place*.

Functions that take a
*const char ** argument are
guaranteeing that they will *not*
modify that argument.

strdup209.c:

```
char *strdup209(const char *s)
{
    size_t len = strlen(s);

    char *dup = (char *) malloc(len + 1);
    if (!dup) {
        return NULL;
    }

    int i;
    for (i = 0; s[i] != '\0'; i++) {
        dup[i] = s[i];
    }
    dup[i] = '\0';

    return dup;
}
```

String Function Summary

- Remember that we are just changing bytes in memory
- *Always* keep in mind who is taking responsibility for the zero/`NUL` terminating byte
 - ... and ultimately, it is still up to you to ensure it is where it is suppose to be!
- Prefer the `str`n— variations whenever possible

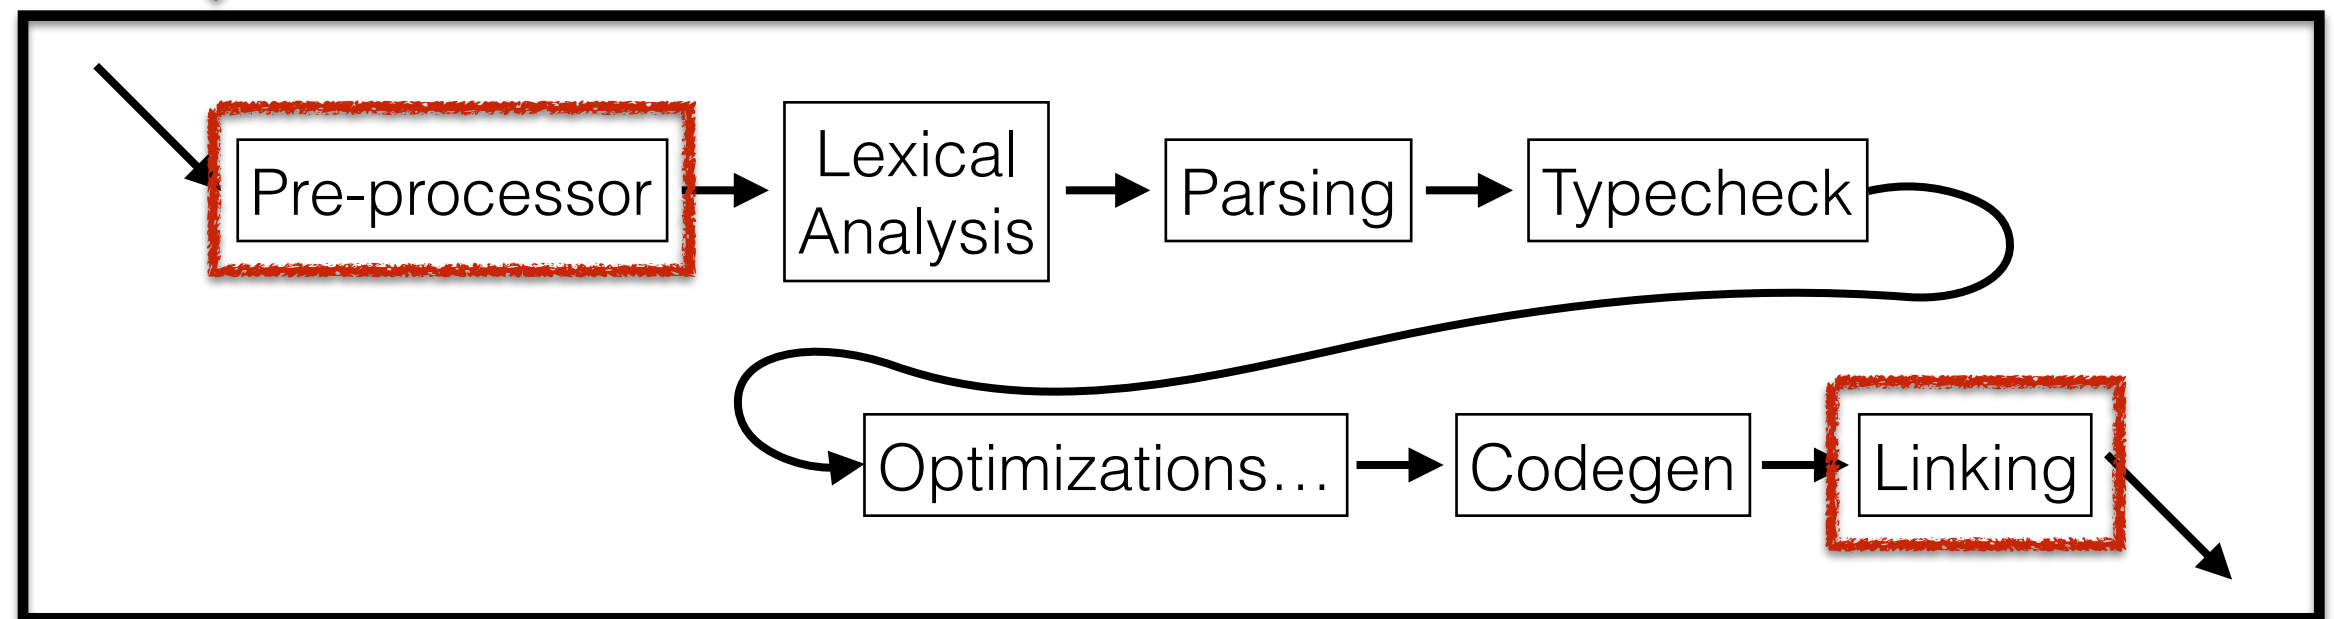
- Find the length of a string:
 - **size_t strlen(const char *s)**
- Compare two strings:
 - **int strcmp(const char *s1, const char *s2)**
 - **int strncmp(const char *s1, const char *s2, size_t n)**
- Find the first/last occurrence of a character within a string
 - **char *strchr(const char *s, int c)**
 - **char *strrchr(const char *s, int c)**
- Append one string onto another:
 - **char *strcat(char *dest, const char *src)**
- Copy one string to another:
 - **char *strncpy(char *dest, const char *src)**
 - Avoid using **char *strcpy(char *dest, const char *src)**
- Duplicate a string:
 - **char *strdup(const char *s)**

Compilation Process

Compilation Process

*.c source code

C Compiler:



Executable binary

Using the C Preprocessor for Sharing Definitions

example.c:

```
//#include <stdlib.h>

...
void *ptr = malloc(...);
...
```

\$ gcc example.c

example.c:15:5: warning: **implicit declaration of function 'malloc'**

Pre-processor Inclusion Mechanism

Original include.c:

```
/* Comments before */  
#include "header.h"  
/* Comments after */
```

header.h:

```
// From `header.h`
```

After pre-processing include.c:

```
/* Comments before */  
// From `header.h`  
/* Comments after */
```

example.c:

```
#include <stdlib.h>

...
void *ptr = malloc(...);
...
```

stdlib.h:

```
...
void *malloc(size_t len);
...
```

After pre-processing example.c:

```
void *malloc(size_t len);

...
void *ptr = malloc(...);
...
```

prog3a.c

```
#include <stdio.h>

extern int flag;

void do_hello();

int main(int argc, char *argv[])
{
    printf("main flag=%d\n", flag);
    flag = 2;

    do_hello();

    printf("main flag=%d\n", flag);

    return 0;
}
```

prog3b.c

```
#include <stdio.h>

int flag = 1;

void do_hello()
{
    printf("hello flag=%d\n",
        flag);
    flag = 3;
}
```

From Assignment 2...

heap209.h:

```
...  
typedef struct _Chunk Chunk; ...  
  
extern void *heap_region;  
extern Chunk *free_list;  
extern Chunk *alloc_list;  
  
void *malloc209(size_t nbytes);  
int free209(void *addr);  
void heap209_init(size_t heap_size);  
void heap209_cleanup(void);  
...
```

Forward declarations of function prototypes and
extern global variables

heap209.c:

```
...  
#include "heap209.h"  
...
```

diagnostics.c:

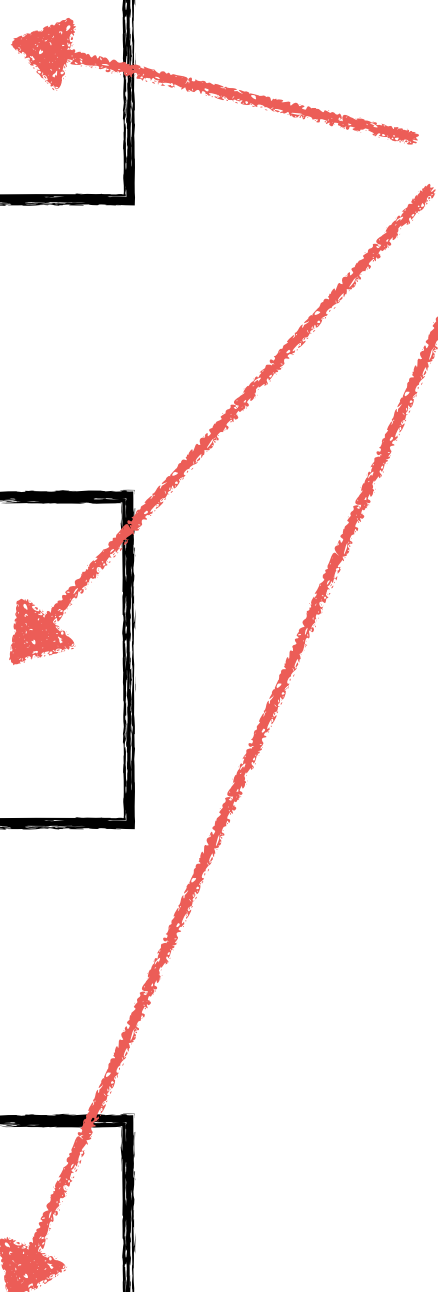
```
...  
#include "heap209.h"  
...
```

test-basic1.c:

```
...  
#include "heap209.h"  
...
```

heap209.h:

```
...  
typedef struct _Chunk Chunk; ...  
  
extern void *heap_region;  
extern Chunk *free_list;  
extern Chunk *alloc_list;  
  
void *malloc209(size_t nbytes);  
int free209(void *addr);  
void heap209_init(size_t heap_size);  
void heap209_cleanup(void);  
...
```



heap209.c:

```
...  
#include "heap209.h"  
...  
  
void *heap_region = NULL;  
Chunk *free_list = NULL;  
Chunk *alloc_list = NULL;  
  
void *malloc209(size_t nbytes)  
{  
    ...  
}  
  
...
```

heap209.h:

```
...  
typedef struct _Chunk Chunk; ...  
  
extern void *heap_region;  
extern Chunk *free_list;  
extern Chunk *alloc_list;  
  
void *malloc209(size_t nbytes);  
int free209(void *addr);  
void heap209_init(size_t heap_size);  
void heap209_cleanup(void);  
...
```

- **heap209.h**

- Forward declarations of function prototypes, **struct** definition and **extern** declarations
- No actual function implementations or global variables defined

- **heap209.c**

- *Implementations* of all the forward declarations in **heap209.h**

- **diagnostics.h**

- Forward declarations of function prototypes for heap debugging functions

- **diagnostics.c**

- *Implementations* of heap debugging functions defined in **diagnostics.h**
- Uses the **struct** definition and **extern** global variables defined in **heap209.h**

- **test-basic1.c**

- Uses the functions defined in **heap209.h** and **diagnostics.h**

What *should* go into header files?

- Function prototypes
- Type definitions: `struct`, `union` and `typedef`
- `extern` global variables
- *Don't declare actual global variables*

badheader.h:

```
int x = 209;
```

file1.c:

```
#include "badheader.h"
```

```
int main()  
{  
    x = 1;  
    return 0;  
}
```

file2.c:

```
#include "badheader.h"
```

```
void file2_utility()  
{  
    x = 2;  
}
```

After pre-processing:

```
int x = 209;  
  
int main()  
{  
    x = 1;  
    return 0;  
}
```

After pre-processing:

```
int x = 209;  
  
void file2_utility()  
{  
    x = 2;  
}
```

```
wolf:~$ gcc -Wall -g file1.c file2.c -o prog
```

```
/tmp/ccdpcL9T.o:(.data+0x0): multiple definition of `x'
```

```
/tmp/ccrZRB1f.o:(.data+0x0): first defined here
```

```
collect2: ld returned 1 exit status
```

Protecting Headers from Multiple Inclusion

- Sometimes the same header file can be included multiple times inadvertently:
 - `prog.c` includes `foo.h`
 - `prog.c` also includes `bar.h`
 - `bar.h` includes `foo.h`
 - `prog.c` will see the contents of `foo.h` twice!

redefine.c:

```
#include "redefine.h"
#include "redefine.h"

int main()
{
    return 0;
}
```

redefine.h:

```
struct S {
    int x;
};
```



After pre-processing:

```
struct S {
    int x;
};
struct S {
    int x;
};

int main()
{
    return 0;
}
```

In file included from redefine.c:2:

./redefine.h:1:8: error: redefinition of 'S'

struct S {
^

./redefine.h:1:8: note: previous definition is here

struct S {
^

1 error generated.

redefine.h:

```
#ifndef REDEFINE_H
#define REDEFINE_H

struct S {
    int x;
};
#endif
```

- On first inclusion:
 - Is REDEFINE_H a currently defined preprocessor symbol?
No (`ifndef` directive is *true*)
 - Define a preprocessor symbol REDEFINE_H
 - Emit body of header file (the `struct` definition)
- On second, third, fourth, etc... inclusion:
 - Is REDEFINE_H a currently defined preprocessor symbol?
Yes (`ifndef` directive is *false*)
 - Skip to `endif` directive

Summary

- Put common definitions into `*.h` (header) files
- Protect your header files from multiple inclusion
- For each function prototype or external global variable, there should be some `*.c` source code file that provides the actual definition/declaration

Summary


- This is one of the core tools for modularity and code reuse when programming in C

Makefiles

Makefiles

- Originally designed to support *separate compilation* of C files

A2 Dependency Graph



Compiling `test-basic1` by hand

```
wolf:~$ gcc -Wall -g -c heap209.c -o heap209.o
```

```
wolf:~$ gcc -Wall -g -c diagnostics.c -o diagnostics.o
```

```
wolf:~$ gcc -Wall -g -c test-basic1.c -o test-basic1.o
```

```
wolf:~$ gcc test-basic1.o heap209.o diagnostics.o \  
        -o test-basic1
```


Anatomy of a Makefile

Target

Prerequisite(s)

Actions(s)

Rule



```
prog : prog.c
```

```
gcc -Wall -g prog.c -o prog
```

- A `Makefile` contains 1 or more *rules*
- Each rule has one *target*, and 1 or more *prerequisites*
- Each rule may have 0 or more *actions* (one per line)

Running make (1)

- `$ make`
 - With no options looks for a file called `Makefile`, and evaluates the *first* rule
- `$ make test-basic2`
 - Looks for a file called `Makefile` and looks for a rule with the target `test-basic2` and evaluates it
- `$ make -f foo ...`
 - Looks for a Makefile with the name *foo*

How it works

- Make looks at the when the target and its prerequisites were **last modified**
 - It assumes targets are files and checks the dates of the files
- Make does nothing...
 - If the target exists, and
 - Is more recent than all its prerequisites
- Make executes the actions...
 - If the target doesn't exist, or
 - If any prerequisite is more recent than the target

Variables — User defined

Define common parts of action commands that you are likely to repeat multiple places:

```
CFLAGS= -Wall -g
```

```
prog : prog.c
```

```
    gcc $(CFLAGS) prog.c -o prog
```

```
prog2 : prog2.c
```

```
    gcc ${CFLAGS} prog.c -o prog
```

Variables — Built-ins

Make defines variables to represent parts of rules:

<code>\$@</code>	Target
<code>\$<</code>	First prerequisite
<code>\$?</code>	All out of date prerequisites
<code>\$^</code>	All prerequisites

```
CFLAGS= -Wall -g
prog : main.c util.c
      gcc ${CFLAGS} $^ -o $@
```

```
CFLAGS= -Wall -g
```

```
prog : main.o util.o
```

```
gcc $^ -o $@
```

```
main.o : main.c
```

```
gcc ${CFLAGS} -c $^ -o $@
```

```
util.o : util.c
```

```
gcc ${CFLAGS} -c $^ -o $@
```

Wildcard Substitutions

```
CFLAGS= -Wall -g
```

```
prog : main.o util.o
```

```
gcc $^ -o $@
```

```
%.o : %.c
```

```
gcc ${CFLAGS} -c $< -o $@
```

A2 Makefile

```
CC=gcc
CFLAGS=-Wall -g
LDFLAGS=
OBJS=heap209.o diagnostics.o

all: test-basic1

test-basic1: test-basic1.o $(OBJS)
    $(CC) $(LDFLAGS) $^ -o $@

%.o: %.c heap209.h
    $(CC) $(CFLAGS) -c $< -o $@

clean:
    rm -f test-basic1 *.o
```

Makefile Summary

- They provide a higher level of abstraction than writing out shell commands in a script file
- They simplify the process of building larger projects

Midterm

- Shell *usage*
- C language:
 - Syntax
 - Data types (including structures and unions)
 - *Pointers and Memory*
- File I/O using streams
- C-style strings
- Makefiles

Midterm

- Be aware of the differences in this course and the courses that previous midterms reflect
- Our midterm date falls in a different week
- We have not yet covered shell *programming*, just emphasized *usage*

- **Labs:** Everyone attending show go to either **BA2210** or **BA2220** (they are conjoined)
- **Extra Office Hours:**
 - Friday, June 19, 1-3pm in BA3289
 - Monday, June 22, 1-3pm in BA3201
- **Midterm:** Tuesday, June 23 from 7-8pm in UC 266
- **Assignment 2:** Due on July 1, 2015