

CSC209 Summer 2015 — Software Tools and Systems Programming

www.cdf.toronto.edu/~csc209h/summer/

Week 8 — July 2, 2015

Peter McCormick
pdm@cs.toronto.edu

Some materials courtesy of Karen Reid

Announcements

- Next Tuesday (July 7) office hour is postponed:
 - Friday, July 10?

Announcements

- *No* lab period tonight
- There is a lab this week (see Piazza) and it is due Sunday 10pm

Agenda

- The course half way point
- Unix processes
- `fork`, `wait` and `exec` system calls

First Half of CSC209

- Shell as user interface
- C language syntax and semantics

Second Half of CSC209

- Mechanisms and abstractions provided by *Unix* operating systems
 - Processes
 - Files
 - Inter-process communication: signals and pipes
 - Network programming with sockets
 - Parallelism and concurrency
- Advanced shell usage

Week 8 lab exercise

Agree *or* Disagree

Agree *or* Disagree

You can only *really* do three things with the C language programming language:

- 1)** Perform arithmetic
- 2)** Evaluate logical expressions
- 3)** Access memory

Unix Processes

Kerrisk ch. 6

“A **process** *is an instance of an executing program.*”

Kerrisk p113

Unix Processes

- Every system has *many* processes currently active:
 - Special operating system processes
 - User process
 - Your personal shell, compiler and test programs...

ps aux

Unix Processes

- Each process is given the illusion of isolation:
 - Exclusive control of the CPU
 - Virtual memory address space
 - Currently open files (including notions of `stdout`, `stdin` and `stderr`)
 - Current working directory
 - Other system resources integral to its execution

Unix Processes

- The OS kernel is an arbitrator that divides physically limited resources out among processes:
 - Memory
 - CPU time
 - Disk
 - Network
 - Access to peripherals, etc.

Unix Processes

- The OS kernel grants each process a slice of the CPU (a short period of time during which the process may run), and then *preemptively* stops that process and switches to let another one run for a time
- Your processes generally do not even notice that this is happening (the illusion of exclusivity)

System Calls

- The kernel lets processes use *system calls* in order to make requests:
 - File I/O
 - Certain kinds of memory management (`mmap`, but not necessarily `malloc`)
 - Process management
 - Communications (networking and IPC)

getpid, getppid - get process identification

```
pid_t getpid(void);  
pid_t getppid(void);
```

From the manpage: “*getpid()* returns the process ID of the calling process.

getppid() returns the process ID of the parent of the calling process.”

getpid.c

What's the difference between a
system library function like

`strncpy`


and a *system call function* like

`getpid`

?

Process State

The **scheduler** decides which of the ready processes to run.



Only one process can be running on a uniprocessor

A process is *ready* if it could use the CPU immediately.

A process is *blocked* if it waiting for an event (I/O, signal)

sleep.c

fork system call

Kerrisk ch. 24

fork

- The `fork` system call creates a copy of the currently running process, diverging from the point of the system call itself:
- The newly created *child* process receives a return value of 0 from the call
- The original *parent* process receives the PID of the newly created child

fork

pid=123:

```
A();  
B();  
C();  
pid = fork();  
// pid == ???  
D();  
E();  
F();
```

fork

pid=123:



```
A( );
```

```
B( );
```

```
C( );
```

```
pid = fork( );
```

```
// pid == ???
```

```
D( );
```

```
E( );
```

```
F( );
```

fork

pid=123:



```
A();
```

```
B();
```

```
C();
```

```
pid = fork();
```

```
// pid == ???
```

```
D();
```

```
E();
```

```
F();
```

fork

pid=123:

A() ;

B() ;

C() ;

pid = fork() ;

// *pid* == ???

D() ;

E() ;

F() ;

fork

pid=123:

```
A();
```

```
B();
```

```
C();
```

```
pid = fork();
```

```
// pid == ???
```


```
D();
```

```
E();
```

```
F();
```


fork

parent (*pid=123*):



```
A();  
B();  
C();  
pid = fork();  
// pid == ???  
D();  
E();  
F();
```


child (*pid=456*):



```
A();  
B();  
C();  
pid = fork();  
// pid == ???  
D();  
E();  
F();
```


fork

parent (*pid=123*):



```
A();  
B();  
C();  
pid = fork();  
// pid == 456  
D();  
E();  
F();
```


child (*pid=456*):



```
A();  
B();  
C();  
pid = fork();  
// pid == ???  
D();  
E();  
F();
```


fork

parent (*pid=123*):



```
A();  
B();  
C();  
pid = fork();  
// pid == 456  
D();  
E();  
F();
```


child (*pid=456*):



```
A();  
B();  
C();  
pid = fork();  
// pid == 0  
D();  
E();  
F();
```



fork

parent (*pid=123*):



```
A();  
B();  
C();  
pid = fork();  
// pid == 456  
D();  
E();  
F();
```


child (*pid=456*):



```
A();  
B();  
C();  
pid = fork();  
// pid == 0  
D();  
E();  
F();
```


fork

parent (*pid=123*):



```
A();  
B();  
C();  
pid = fork();  
// pid == 456  
D();  
E();  
F();
```


child (*pid=456*):



```
A();  
B();  
C();  
pid = fork();  
// pid == 0  
D();  
E();  
F();
```


fork

parent (*pid=123*):



```
A();  
B();  
C();  
pid = fork();  
// pid == 456  
D();  
E();  
F();
```

child (*pid=456*):




```
A();  
B();  
C();  
pid = fork();  
// pid == 0  
D();  
E();  
F();
```

fork


parent (*pid=123*):

```
A();  
B();  
C();  
pid = fork();  
// pid == 456  
D();  
E();  
F();
```




child (*pid=456*):

```
A();  
B();  
C();  
pid = fork();  
// pid == 0  
D();  
E();  
F();
```




fork

parent (*pid=123*):



```
A();  
B();  
C();  
pid = fork();  
// pid == 456  
D();  
E();  
F();
```

child (*pid=456*):



```
A();  
B();  
C();  
pid = fork();  
// pid == 0  
D();  
E();  
F();
```

fork

parent (*pid=123*):

```
A();  
B();  
C();  
pid = fork();  
// pid == 456  
D();  
E();  
F();
```

child (*pid=456*):

```
A();  
B();  
C();  
pid = fork();  
// pid == 0  
D();  
E();  
F();
```

fork

parent (*pid=123*):

```
A();  
B();  
C();  
pid = fork();  
// pid == 456  
D();  
E();  
F();
```

child (*pid=456*):

```
A();  
B();  
C();  
pid = fork();  
// pid == 0  
D();  
E();  
F();
```



fork

parent (*pid=123*):

```
A();  
B();  
C();  
pid = fork();  
// pid == 456  
D();  
E();  
F();
```

child (*pid=456*):

```
A();  
B();  
C();  
pid = fork();  
// pid == 0  
D();  
E();  
F();
```



fork

parent (*pid=123*):

```
A();  
B();  
C();  
pid = fork();  
// pid == 456  
D();  
E();  
F();
```

child (*pid=456*):

```
A();  
B();  
C();  
pid = fork();  
// pid == 0  
D();  
E();  
F();
```

fork.c

fork

- Fork can and will *fail* (i.e. return `-1`) if your user account has created too many processes, or if a system wide limit has been reached

`fork` — what's initially the same between parent and child?

- Properties of parent *inherited* by child:
 - UID, GID
 - Controlling terminal
 - Current working directory and notion of root directory
 - Signal mask, environment, resource limits
 - Shared memory (SHM) segments

`fork` — what's changes between parent and child?

- *Differences* between parent and child
 - PID, PPID
 - *Return value* from `fork ()`
 - Pending *alarms* cleared for child
 - Pending *signals* are cleared for child

Process Termination

Kerrisk ch. 25

Process Termination

- A process *terminates* when either it explicitly calls `exit(int status)` or implicitly when it *returns* from `main` with a status code
 - Status code of 0 indicates *success* (or, the absence of failure)
 - Anything else indicates *failure*
- The Bash shell stores the exit status code of the last process run in a special variable named `$?`

exitstatus.c

Process Termination

- Every normal process is the child process of some parent process
- A terminating process sends its parent a **SIGCHLD** signal and waits for its parent to accept its exit code

What happens if the
parent *exits* before the
child?

Orphaned Processes

- Any process whose parent terminates before it does will become *orphaned*, and its parent process becomes PID 1 (the `init` process, which is the first process in the entire system)

forkorphan.c

How does a parent process
wait for its child to exit
before itself terminating?

`wait` - wait for child process to change state

```
pid_t wait(int *status);
```

- A process that calls `wait()` can:
 - *block* (if all of its children are still running)
 - return immediately with the termination status of a child (if a child has terminated and is waiting for its termination status to be fetched)
 - return immediately with an error (if it doesn't have any child processes.)

wait.c

Zombies

- *A zombie process:*
 - a process that is “waiting” for its parent to accept its exit status code
 - a parent accepts a child’s status code by executing `wait()`
 - shows up as `Z` in `ps -a`
 - A terminating process may be a (multiple) parent; the kernel ensures all of its children are orphaned and adopted by `init`

zombie.c

wait and waitpid

- `wait()` can
 - block
 - return with termination status
 - return with error
- If there is more than one child, `wait()` returns on termination of *any* of its children
- `waitpid()` can be used to wait for a *specific* child PID
 - Also has an option to block or not to block

```
pid_t waitpid(pid_t pid,  
               int *status,  
               int options);
```

- if `pid == -1`:
 - Wait for *any* child (otherwise wait only for that specific child `pid`)
- if `option == WNOHANG`:
 - Return immediately *if* there is no child to wait for (i.e. do *not* block)
- if `option == 0`:
 - *Do* wait (block) until there is a child to deal with

`wait(&status)` is equivalent to `waitpid(-1, &status, 0)`

waitpid.c

waitmany.c
and
kill(1)

`fork` is the only way to
create new processes

... how do we ever run
existing programs then?

exec

Kerrisk ch. 27

`exec` - replace the currently running process

- A family of system calls with several different variations
- Replaces the program that the process is currently running with another
- On success, `exec` will *never* return (because success means another program is now running in your place), and on failure will return `-1`

./example (pid=**123**):

...



`exec*("/bin/ls");`

// Never run...

./example (pid=**123**):

```
...  
exec*("/bin/ls");  
// Never run...
```



/bin/ls (pid=**123**):

```
...  
code for /bin/ls  
...
```

Properties of `exec`

- New process *inherits* from calling process:
 - *PID and PPID*
 - Real UID, GID
 - Controlling terminal
 - CWD, root directory, resource limits
 - Pending signals
 - Pending alarms

Variations of exec

```
int execl(const char *filename, char *const argv[], char *const envp[]);  
int execlp(const char *path, char *const argv[]);
```

```
int execl(const char *path, const char *arg, ...);  
int execlp(const char *path, const char *arg, ..., char *const envp[]);  
int execlp(const char *file, const char *arg, ...);
```

```
int execlp(const char *file, char *const argv[]);  
int execlpe(const char *file, char *const argv[],  
            char *const envp[]);
```

Variations: `execv`

```
int execv(const char *path,  
            char *const argv[]);
```

Exec the binary executable located at `path`, passing in the given `argv` array (which must be `NULL` terminated)

`execv.c`

Variations: `exec*``p``*`

```
int execlp(const char *file, char *const argv[]);  
int execlp(const char *file, const char *arg, ...);  
int execvp(const char *file, char *const argv[],  
           char *const envp[]);
```

Use the `PATH` environment variable to search for executables with the name specified in `file`

execvp.c

Variations: `exec*1*`

```
int exec1(const char *path, const char *arg, ...);  
int exec1p(const char *file, const char *arg, ...);  
int exec1e(const char *path, const char *arg, ...,  
           char * const envp[]);
```

Uses C *variadic* functions (which allow a variable number of parameters) to let you specify the contents of `argv` (must have signal the end with an explicit `NULL`.)

Variations: `exec*e`

```
int execve(const char *filename, char *const argv[],  
            char *const envp[]);
```

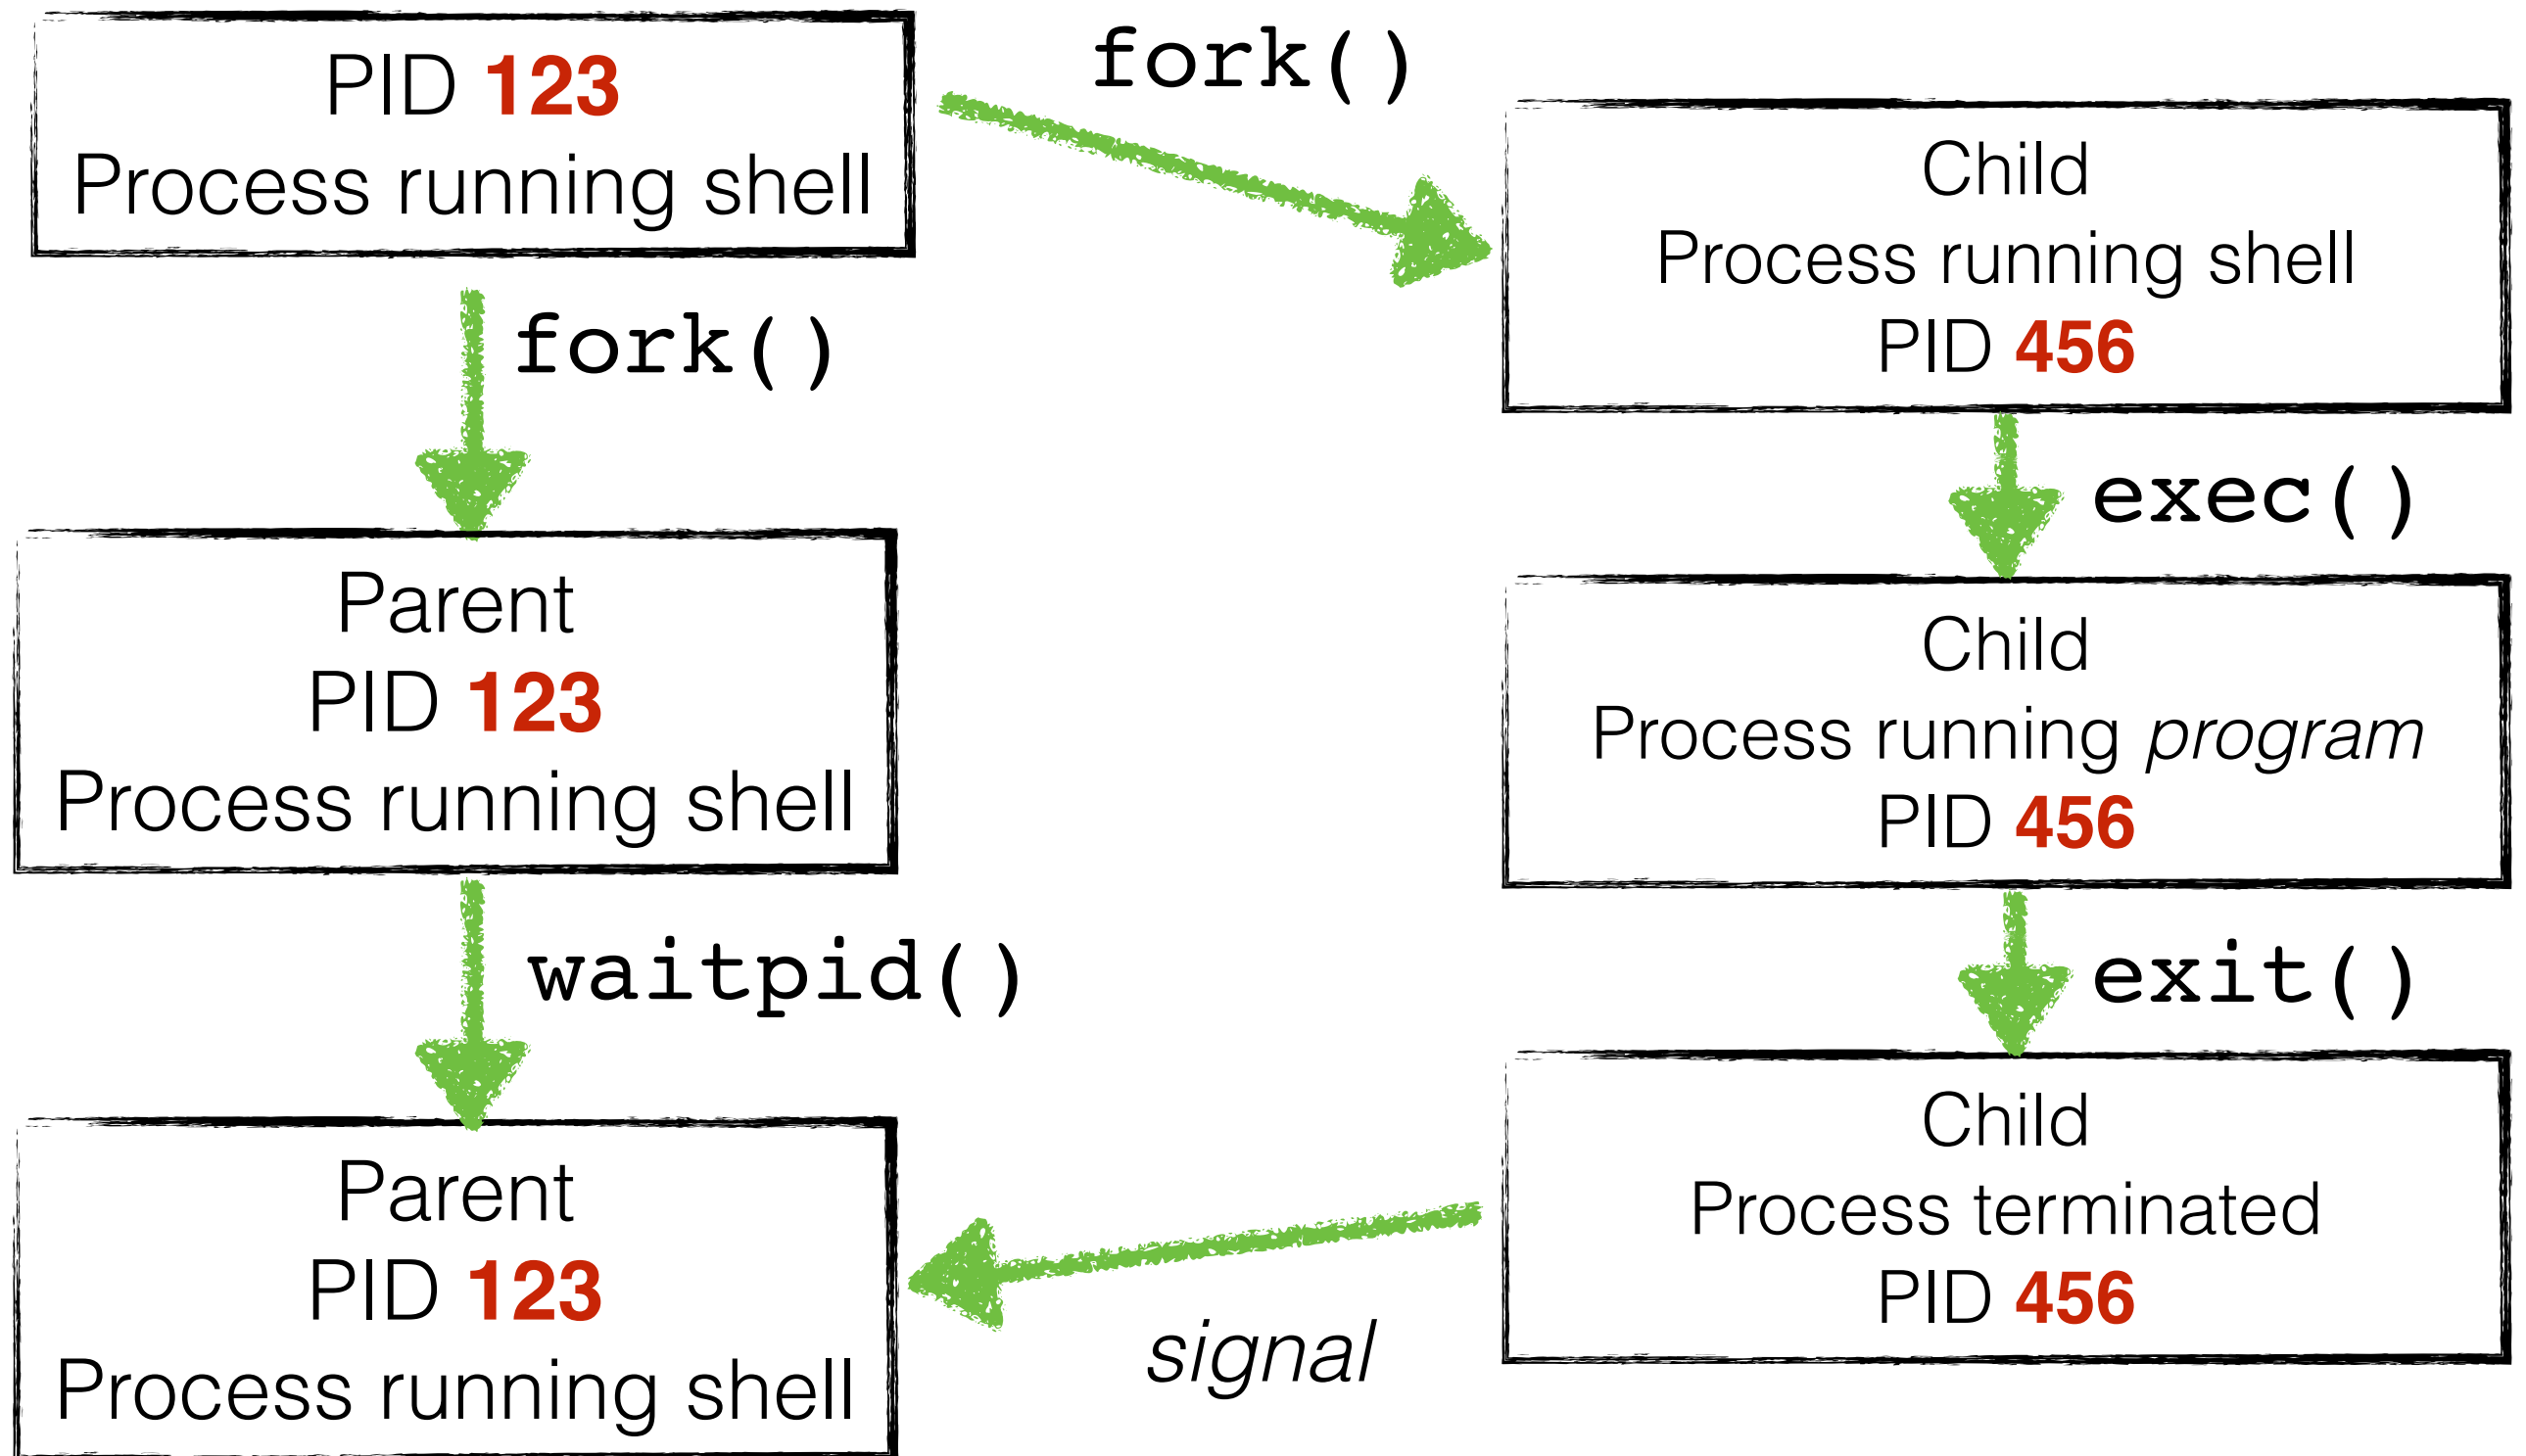
```
int execvpe(const char *file, char *const argv[],  
            char *const envp[]);
```

```
int execle(const char *path, const char *arg, ...,  
            char * const envp[]);
```

Specify the *environment* (`envp` array) to `exec` the program in.

forkexec.c

How a shell works



`fork` is the only way to
create new processes

`exec` is the only way to
run existing programs

Midterms